Lab 4: Overlay Network and VXLAN

518030910283 王航宇

First, I get the two VMs' IP address by the command:

```
ifconfig
```

IP address on VM1: 192.168.164.3

IP address on VM2: 192.168.164.4

Then I construct the network topology on VM1 and VM2 by the files "build_1.py" and "build_2.py". All the links bandwith is 10Mbps, the packet loss rate is 0%, and the time delay is 5ms

```
class MultiSwitchTopo(Topo):
    "Multi switches connected to hosts."

    def build(self):
        s1 = self.addSwitch('s1')
        h1 = self.addHost('h1')
        h2 = self.addHost('h2')
        self.addLink(s1, h1,bw=10,loss=0,delay='5ms')
        self.addLink(s1, h2,bw=10,loss=0,delay='5ms')

def simpleTest():
    "Create and test a simple network"
    topo = MultiSwitchTopo()
    net = Mininet(topo, link=TCLink)
    net.start()
    CLI(net)
    net.stop()
```

Then I follow the steps in the file lab4-questions.pdf. Note that I use 192.168.164.5 as the default route for br1:

```
sudo route add default gw 192.168.164.5
```

Homework 1:

(30 points) Ping 10.0.0.111 from 10.0.0.112 (in your terminal of VM2) and use Wireshark to monitor the interfaces s2 and enp0s8, and describe the protocols used in this procedure and your findings.

Requirements:

Please provide the screenshot of the wireshark interface in your report. Your description and findings should align with the information in wireshark interface

In the terminal of VM2, I enter the command to ping 10.0.0.111:

```
ping 10.0.0.111
```

The screenshot of Wireshark monitoring the interfaces s2 is:

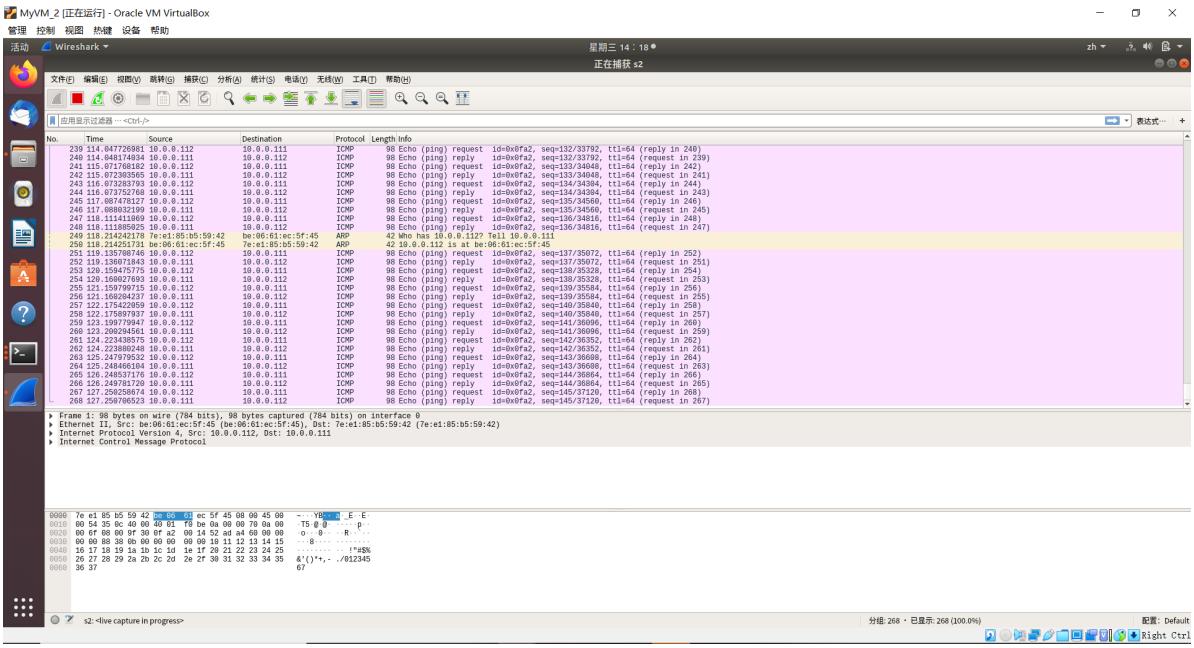

From this screenshot, we can see that Internet Control Message Protocol(ICMP), Internet Protocol Version 4(IPv4) and Address Resolution Protocol(ARP) are used in this procedure.

The screenshot of Wireshark monitoring the interfaces enp0s8 is:

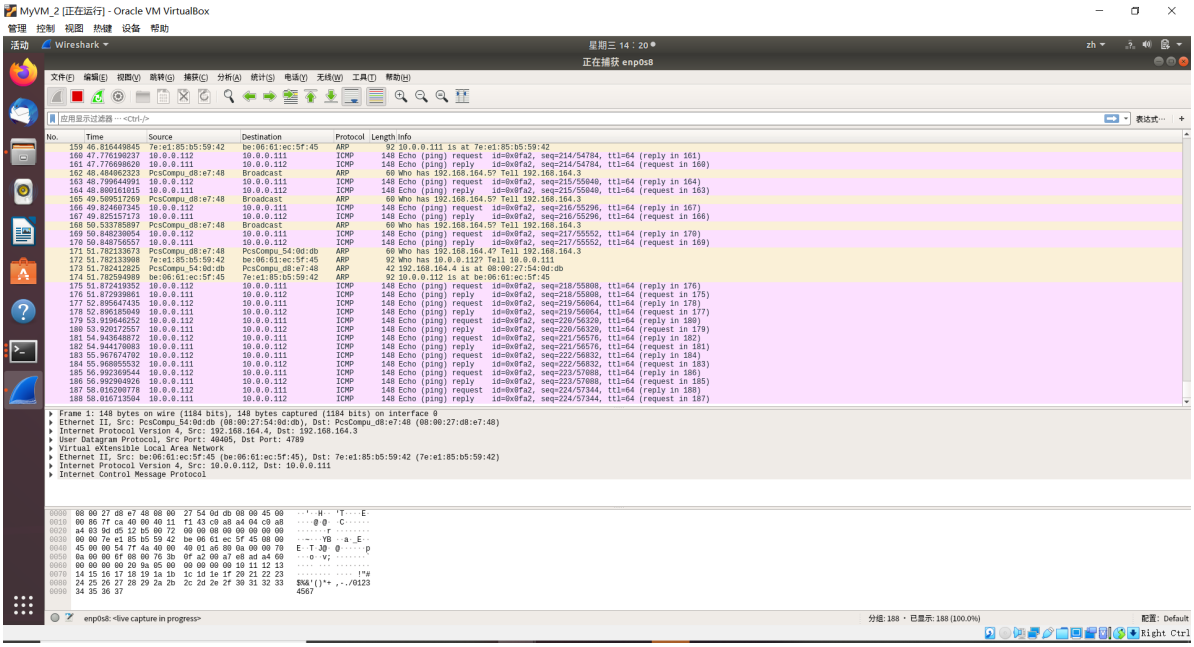

From this screenshot, we can see that Internet Control Message Protocol(ICMP), Internet Protocol Version 4(IPv4), Address Resolution Protocol(ARP) and User Datagram Protocol(UDP) are used in this procedure.

Decriptions:

ICMP is a network layer protocol used by network devices to diagnose network communication issues. ICMP is mainly used to determine whether or not data is reaching its intended destination in a timely manner.

IPv4 is a connectionless protocol, and operates on a best-effort delivery model, in that it does not guarantee delivery, nor does it assure proper sequencing or avoidance of duplicate delivery.

ARP is a protocol used by the Internet Protocol, specifically IPv4, to map IP network addresses to the hardware addresses used by a data link protocol.

UDP is a lightweight data transport protocol that works on top of IP. UDP provides a mechanism to detect corrupt data in packets, but it does not attempt to solve other problems that arise with packets, such as lost or out of order packets.

Findings:

I find that enp0s8 and s2 both use ICMP, IPv4 and ARP, but enp0s8 uses one more protocol than s2, that's UDP. Moreover, the frame length in enp0s8 is 148 bytes, which is longer than 98 bytes, the frame length in s2. The extra 50 bytes are added by VXLAN.

Homework 2:

(50 points) Use iperf to test the network bandwidth between the two virtual machine

• Test the bandwidth between 192.168.56.101 and 192.168.56.102

• Test the bandwidth between 10.0.0.1/10.0.0.2/10.0.0.111 and 10.0.0.3

Compare the above results and explain the reason. (Hint: you may need to specify a reasonable MTU size in order for your iperf to work in this case. Please also think about why.)

Requirements:

Please provide the screenshot of your terminals in your report.

In the terminal of VM1, I enter the command to test the bandwidth between 192.168.164.4 and 192.168.164.3:

```
iperf3 -t 10 -i 1 -c 192.168.164.4 -p 16000
```

and in VM2, I enter:

```
iperf3 -s -i 1 -p 16000
```

The bandwidth between 192.168.164.3 and 192.168.164.4 is 2.16 Gbits/sec

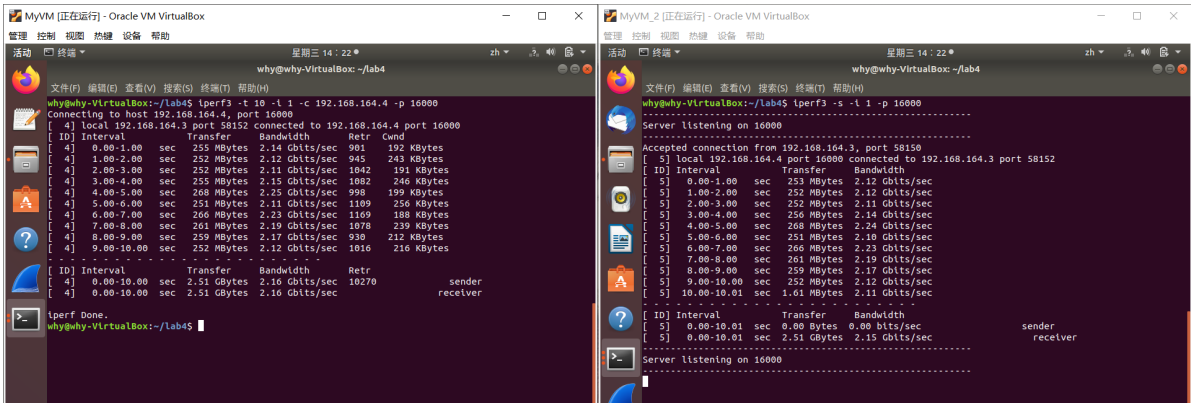

In the xterm of 10.0.0.1/10.0.0.2/10.0.0.111, I all enter:

```
iperf3 -M 1410 -t 10 -i 1 -c 10.0.0.3 -p 16000
```

and in the xterm of 10.0.0.3, I enter:

```
iperf3 -s -i 1 -p 16000
```

-M set the maximum segment size(= $MTU - 40$ bytes), I find the maximum value of -M is 1410 (so MTU$= 1450$) and if -M is 1411, the iperf won't work. That's because the default MTU of the virtual machine is 1500 Bytes, and when this message passes through the VTEP (VXLAN Tunnel End Point), a new header of 50 bytes will be encapsulated. so we need to lower MTU size(maximum MTU=1500-50=1450) in order to avoid the segment size exceeding the limit.

The bandwidth between 10.0.0.1 and 10.0.0.3 are 9.49 Mbits/sec.

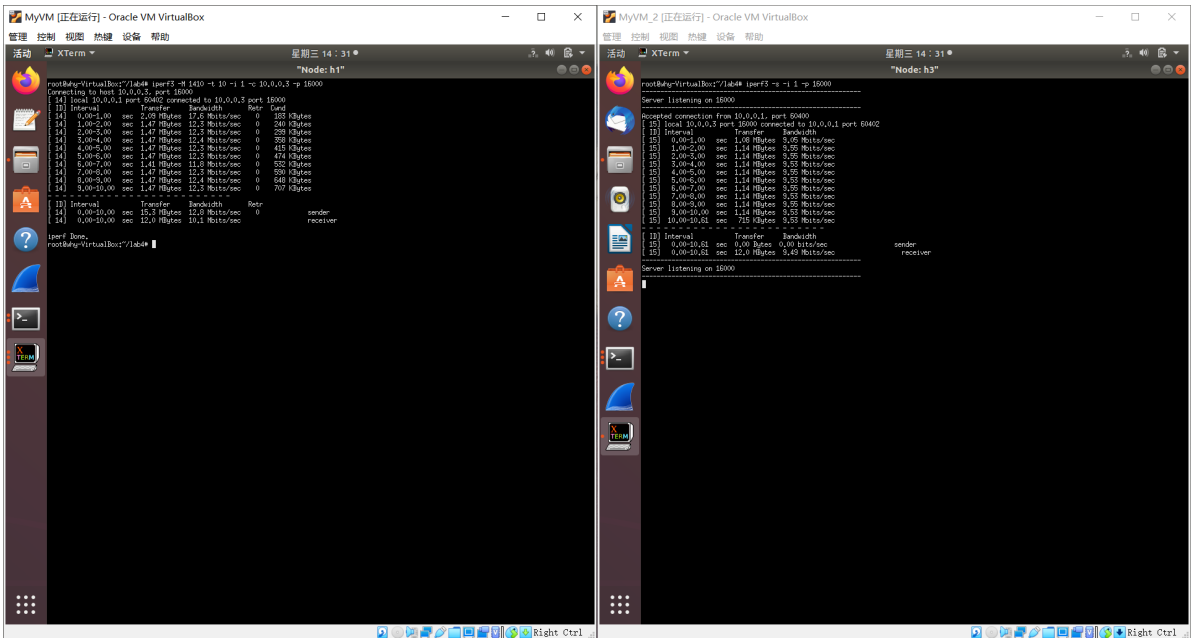

The bandwidth between 10.0.0.2 and 10.0.0.3 are 9.50 Mbits/sec.

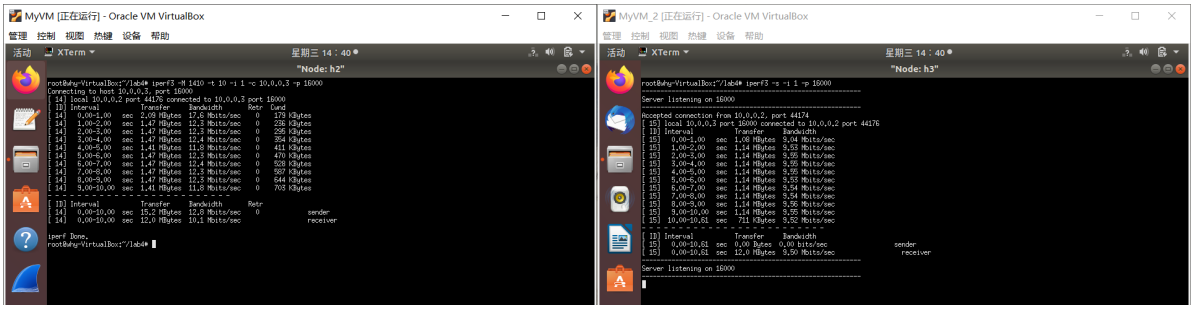

The bandwidth between 10.0.0.111 and 10.0.0.3 are 9.53 Mbits/sec.

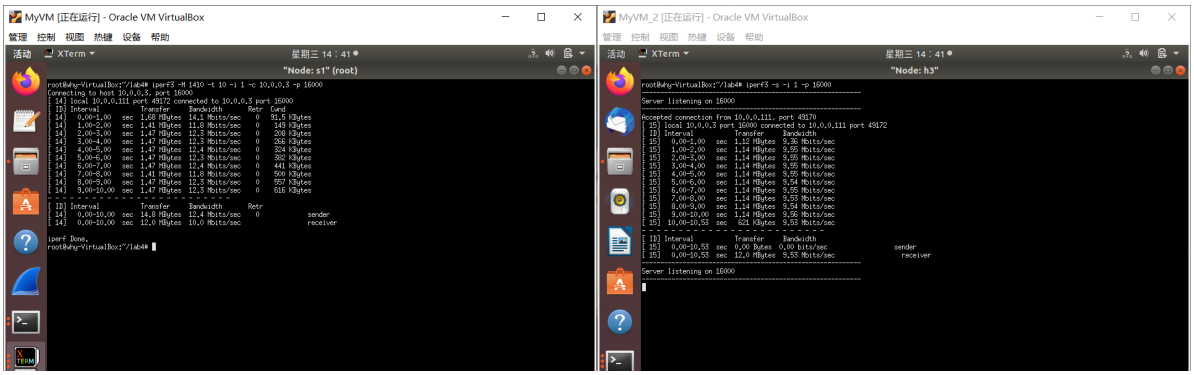

We can find that

- the bandwidths between 10.0.0.1/10.0.0.2/10.0.0.111 and 10.0.0.3 are almost the same
- the bandwidth between 192.168.164.4 and 192.168.164.3 is larger than the bandwidths between 10.0.0.1/10.0.0.2/10.0.0.111 and 10.0.0.3.

That's because we set all the links bandwidth as 10 Mbps, so even the real bandwidth can reach 2.16 Gbits/sec, the bandwidth in the network topology can't exceed 10 Mbps.

Homework 3:

(20 points) Similar to Q2, use ping to test the network latency and analyze your results.

Requirements:

Please provide the screenshot of your terminals in your report.

In the terminal of VM1, I enter the command to test the RTT between 192.168.164.4 and 192.168.164.3:

```
ping -c 20 192.168.164.4
```

The average RTT between 192.168.164.4 and 192.168.164.3 is 0.487 ms.

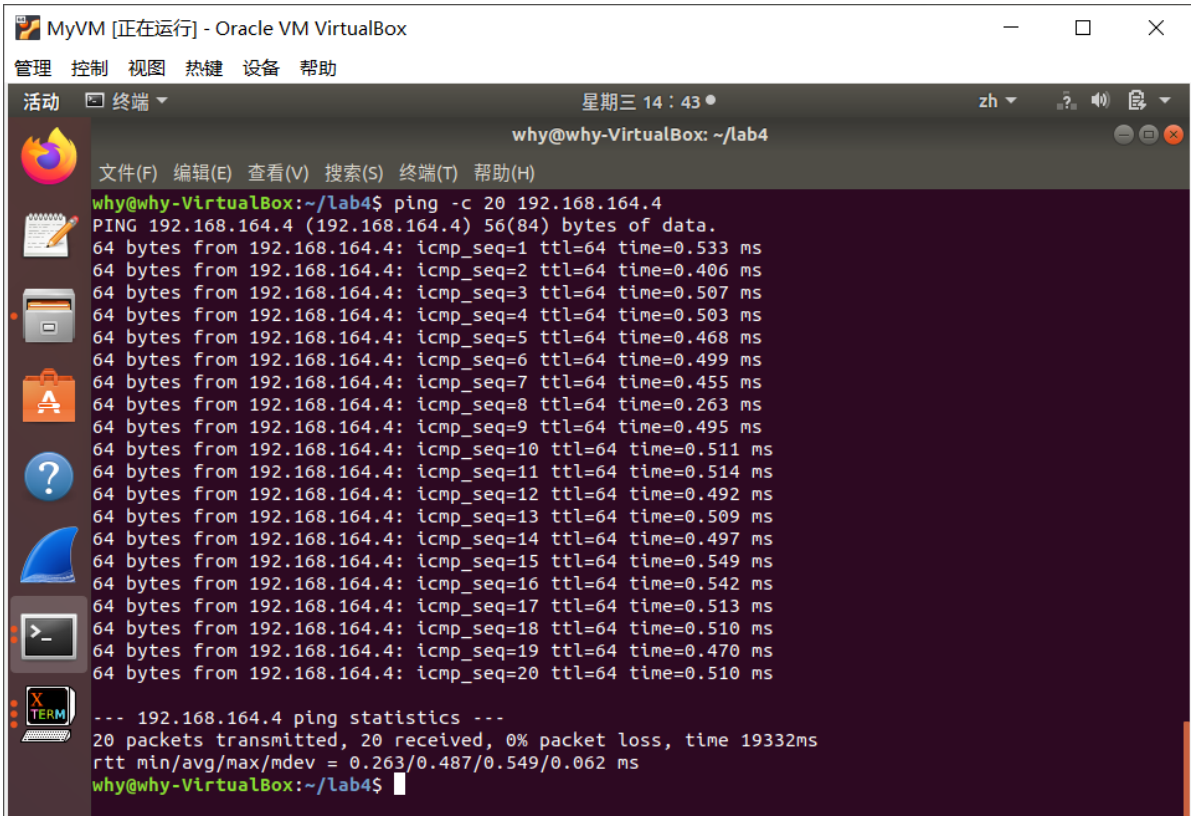

Then I test the latency between 10.0.0.1/10.0.0.2/10.0.0.111 and 10.0.0.3 respectively:

```
ping -c 20 10.0.0.3
```

The average RTT between 10.0.0.1 and 10.0.0.3 is 22.643 ms.

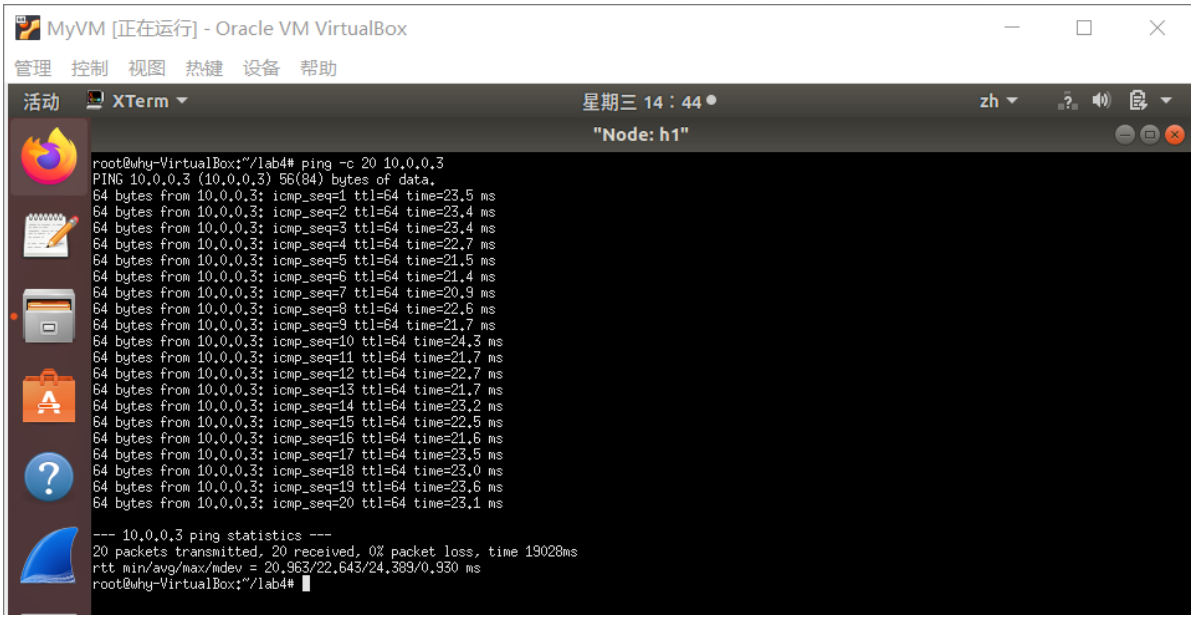

The average RTT between 10.0.0.2 and 10.0.0.3 is 22.777 ms.

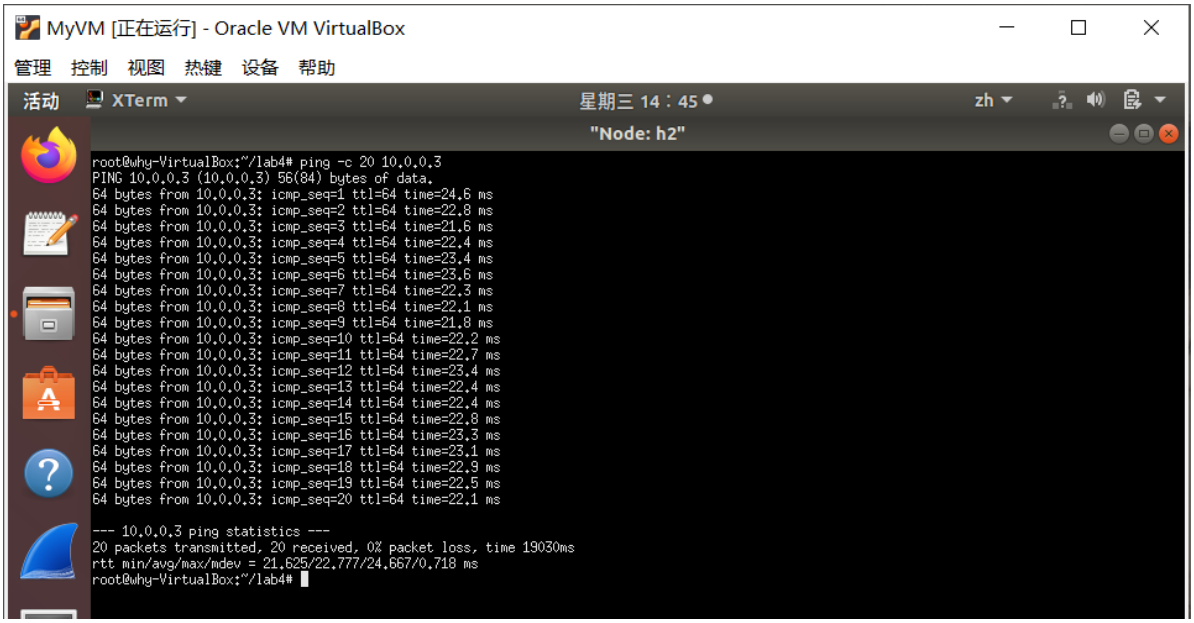

The average RTT between 10.0.0.111 and 10.0.0.3 is 11.525 ms.

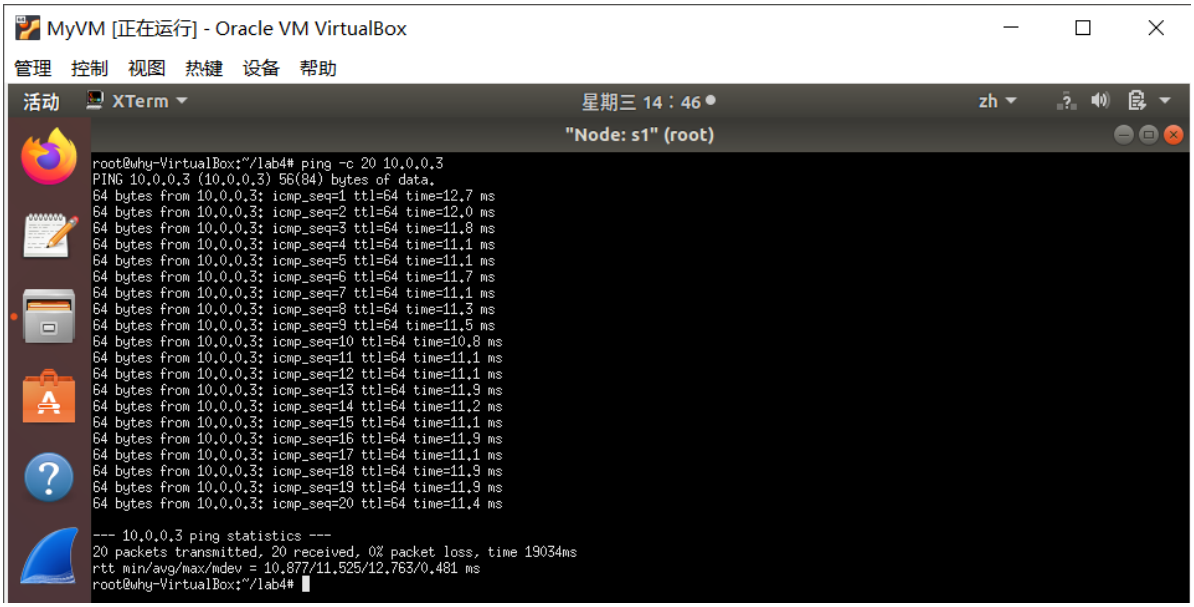

We can find that

1. The average RTTs between 10.0.0.1/10.0.0.2 and 10.0.0.3 are almost the same (22 ms)
2. The average RTTs between 10.0.0.1/10.0.0.2 and 10.0.0.3 are larger than the average RTT between 10.0.0.111 and 10.0.0.3 (11 ms)
3. The average RTTs between 192.168.164.4 and 192.168.164.3 is smallest (0.487ms)

(1) is because the topological status of h1 and h2 are equivalent, so their latency are almost the same

(2) is because we set all the links time delay is 5 ms, and a round trip from h1 (or h2) to h3 passes links 4 times (4*5=20ms), a round trip from s1 to h3 passes links 2 times (2*5=10ms)

(3) 's reason is obvious, it shows that the real latency is very small.